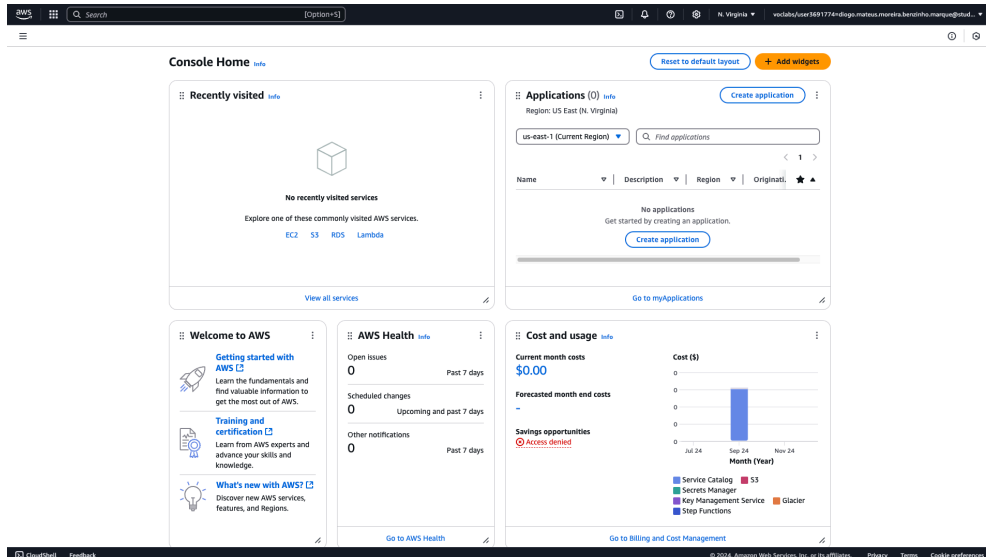



# LS Lab Assignment: Amazon EC2\*

## 1 Access



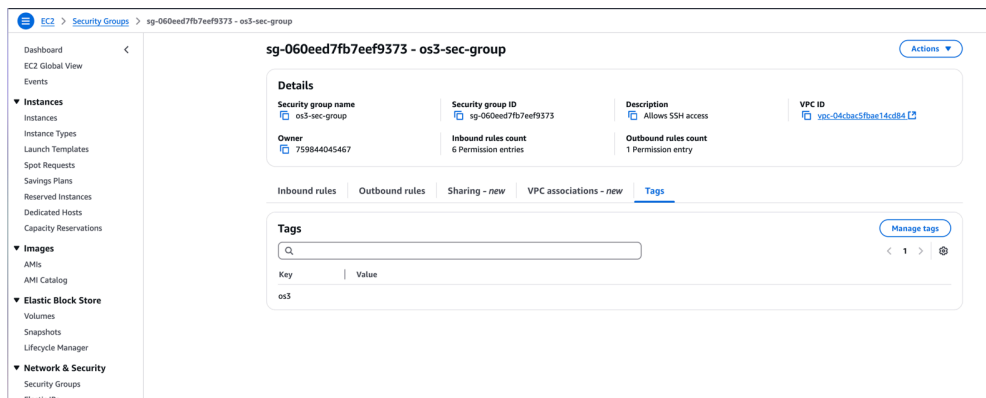
dashboard

## 1 Key pair os3



key-pair.png

## 2 Network security group



network-security-group.png

## 2 Adding Instances

### 3 What is the difference between EBS and instance-store for VMs?

The key difference between **EBS** (Elastic Block Store) and **instance-store** lies in data persistence and use cases.

**EBS** is a persistent, durable block storage service that remains intact even when the EC2 instance is stopped or terminated. It is separate from the instance, allowing it to be detached and reattached to other instances. **EBS** supports snapshots, encryption, and offers customizable storage options, making it ideal for applications requiring long-term data storage, such as databases or file systems.

In contrast, **instance-store** is ephemeral storage physically attached to the host machine of the EC2 instance. Data in **instance-store** is temporary and lost when the instance is stopped or terminated. It does not support snapshots or customization and is typically used for temporary tasks like caching or data buffering that do not require persistence.

Feature	EBS	Instance-Store
<b>Persistence</b>	Persistent	Ephemeral
<b>Attached to Instance</b>	Detachable and re-attachable	Fixed to the instance
<b>Data Loss on Stop/Restart</b>	No	Yes
<b>Scalability</b>	Adjustable size	Fixed size based on instance type
<b>Performance</b>	Suitable for varied workloads	High-speed, local access
<b>Snapshot Support</b>	Yes	No
<b>Backup Options</b>	AWS Snapshots	None
<b>Use Cases</b>	Databases, application storage	Caching, temporary processing
<b>Cost</b>	Separate charge for volumes	Included in EC2 cost

In summary, **EBS** is suited for workloads requiring persistent and scalable storage, while **instance-store** is optimal for high-speed, temporary storage needs.

## 4 Add running instances to security group

### Security network group inbound rules

The screenshot shows the AWS Management Console for the security group **sg-060eed7fb7eef9373 - os3-sec-group**. A green notification banner at the top states: "Inbound security group rules successfully modified on security group (sg-060eed7fb7eef9373 | os3-sec-group)".

**Details:**

- Security group name: os3-sec-group
- Security group ID: sg-060eed7fb7eef9373
- Description: Allows SSH access
- VPC ID: vpc-04cbac5fbac14cd84
- Owner: 759844045467
- Inbound rules count: 2 Permission entries
- Outbound rules count: 1 Permission entry

**Inbound rules (2):**

Name	Security group rule ID	IP version	Type	Protocol	Port range
ssh access	sgr-0135ac643d6422f88	IPv4	SSH	TCP	22
http access	sgr-0714a97ef7eda2803	IPv4	HTTP	TCP	80

### Security network group outbound rules

The screenshot shows the AWS Management Console for the security group **sg-060eed7fb7eef9373 - os3-sec-group**. The **Outbound rules** tab is selected.

**Details:**

- Security group name: os3-sec-group
- Security group ID: sg-060eed7fb7eef9373
- Description: Allows SSH access
- VPC ID: vpc-04cbac5fbac14cd84
- Owner: 759844045467
- Inbound rules count: 1 Permission entry
- Outbound rules count: 1 Permission entry

**Outbound rules (1):**

Security group rule ID	IP version	Type	Protocol	Port range	Destination
sgr-05334c7e256d31e17	IPv4	All TCP	TCP	0 - 65535	0.0.0.0/0

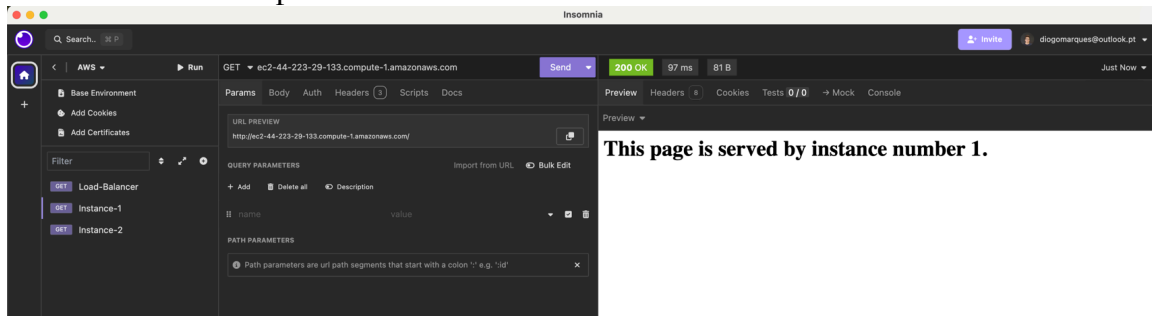
## 5 Create and Test the WebServers

### Commands used to install and server the webserver for Linux machines.

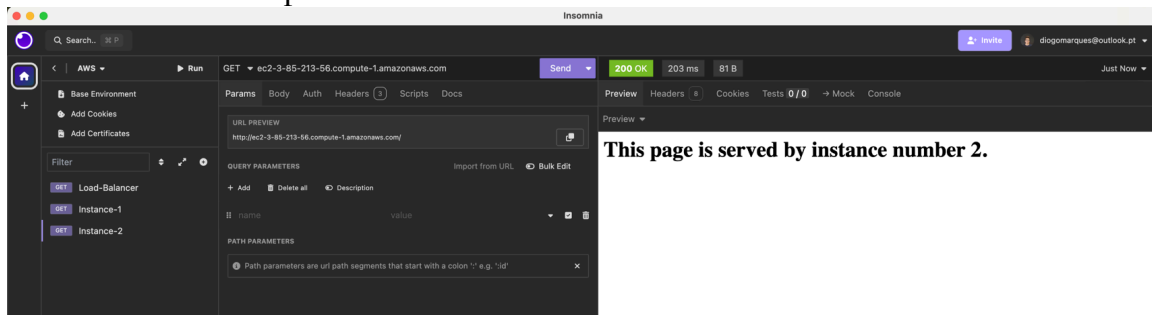
```
sudo apt-get update
sudo apt-get install apache2 -y
sudo nano /var/www/html/index.html
cat /var/www/html/index.html
<html>
<body>
<h1>This page is served by instance number #.</h1>
</body>
</html>
```

## 6, 7 Request test done with AWS assigned DNS and measure response time

### Instance 1. 97ms response time



### Instance 2 203ms response time



Create the script `measure_response_time.py` to perform the http requests.

```
import httpx
import asyncio
import numpy as np
import time
```

*# Function to calculate metrics from a list of response times*

```
def calculate_metrics(response_times):
    """
```

*Calculates statistical metrics for a list of response times.*

*Args:*

*response\_times (list): List of response times in seconds.*

*Returns:*

*dict: A dictionary containing metrics: average, median, minimum, maximum, and standard deviation.*

```
    """
    return {
        "Average": np.mean(response_times),
        "Median": np.median(response_times),
        "Minimum": np.min(response_times),
        "Maximum": np.max(response_times),
        "Standard Deviation": np.std(response_times)
    }
```

*# Asynchronous function to measure response times for multiple URLs n times*

```
async def measure_multiple_async_response_times(urls, n):
```

```
"""
```

*Measures response times for multiple URLs asynchronously, repeated n times.*

*Args:*

*urls (list): List of URLs to send GET requests to.*  
*n (int): Number of times to repeat the requests.*

*Returns:*

*list: A list of response times in seconds for all requests.*

```
"""
```

```
async with httpx.AsyncClient() as client:
    response_times = []

    for i in range(n):
        tasks = [client.get(url) for url in urls] # Create a task
        for each URL
        start_time = time.time()
        responses = await asyncio.gather(*tasks) # Wait for all
        tasks to complete
        end_time = time.time()

        # Calculate response times for this batch of URLs
        batch_time = end_time - start_time
        response_times.append(batch_time)
        print(f"Batch {i + 1}: Response time: {batch_time:.4f}
        seconds")

    return response_times
```

*# Main function to demonstrate asynchronous response time measurement*

```
if __name__ == "__main__":
```

```
    """
```

*Main function to measure response times for URLs run n times and calculate metrics.*

```
    """
```

*# List of URLs to test*

```
urls = [
    "http://ec2-44-223-29-133.compute-1.amazonaws.com",
    # Ubuntu instance - 1 DNS
    "http://54.164.184.234",
    # Ubuntu instance - 1 IP
    "http://ec2-3-85-213-56.compute-1.amazonaws.com",
    # Ubuntu instance - 2 DNS
    "http://34.207.182.80",
    # Ubuntu instance - 2 IP
    "http://ec2-34-207-147-155.compute-1.amazonaws.com",
    # Windows instance
    "http://os3-load-balancer-1369570309.us-east-1.elb.amazonaws.com" # Load Balancer
]
```

*# Number of times to repeat the requests*

```
n = 50
```

*# Measure response times asynchronously*

```
print("Asynchronous Response Time Measurement:")
```

```

async_response_times =
    asyncio.run(measure_multiple_async_response_times(urls, n))

# Calculate and print metrics for asynchronous requests
async_metrics = calculate_metrics(async_response_times)
print(f"\nAsynchronous Response Time Metrics for {n} Runs:")
for metric, value in async_metrics.items():
    print(f"{metric}: {value:.4f} seconds")

```

There are the results for 50 requests to the Ubuntu Web server instances.

Asynchronous Response Time Metrics for 50 Runs:

Average: 0.1007 seconds

Median: 0.0983 seconds

Minimum: 0.0968 seconds

Maximum: 0.1988 seconds

Standard Deviation: 0.0142 seconds

## Results for 50 requests.

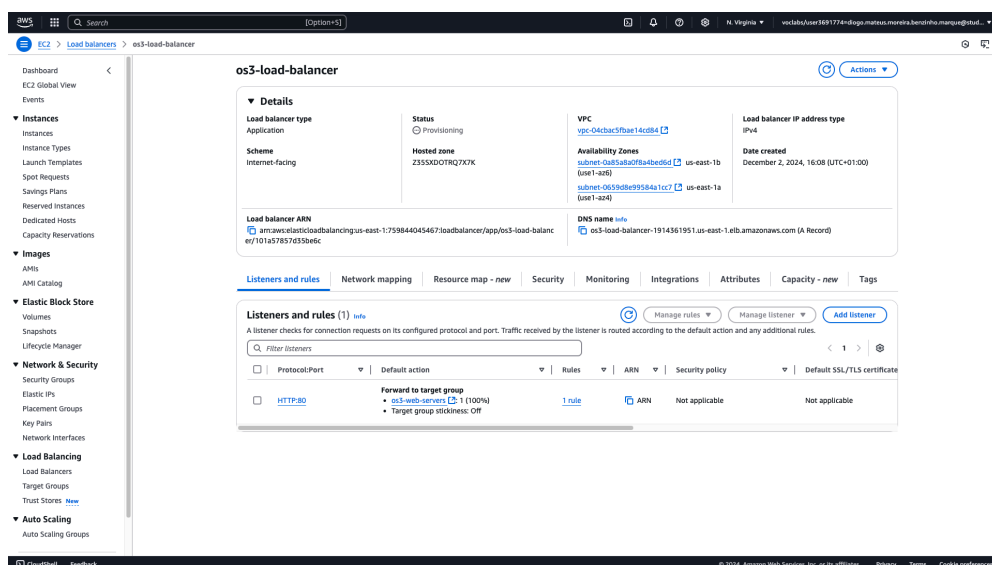
**Metrics:** - Average Response Time: 0.1007 seconds - Median: 0.0983 seconds - Minimum: 0.0968 seconds - Maximum: 0.1988 seconds Standard Deviation: 0.0142 seconds

## Observations:

- The response times are highly consistent across batches, with the majority of batches averaging around 0.098 seconds.
- A single outlier in Batch 1 (0.1988 seconds) significantly increases the maximum response time.
- The low standard deviation indicates stable server performance without much variability between requests.

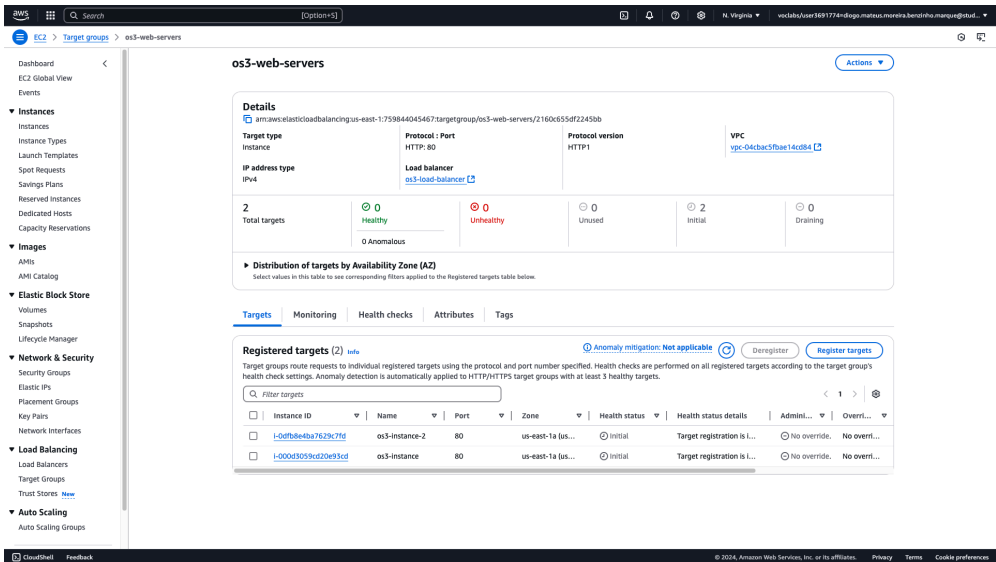
# 3 Optional: Load Balancing

## Creation of Load Balancer



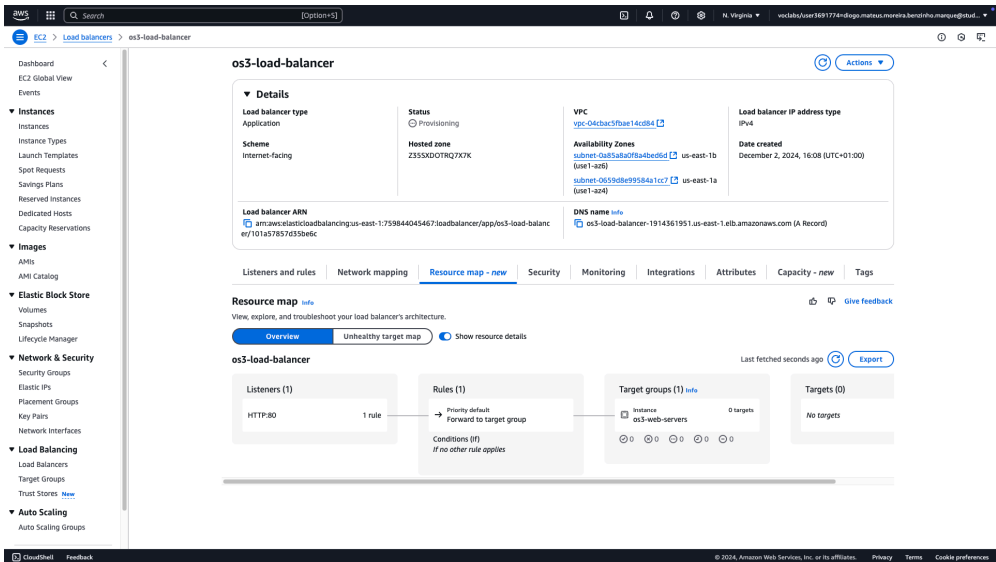
load-balancer.png

## Target groups



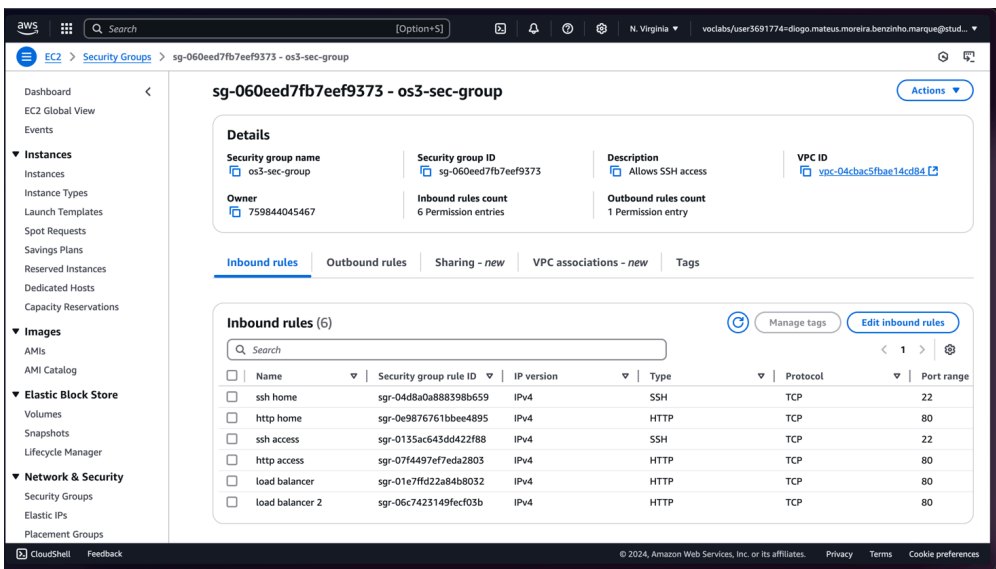
load-balancer-target-group.png

## Confirmation with Resource Map view tool



load-balancer-2.png

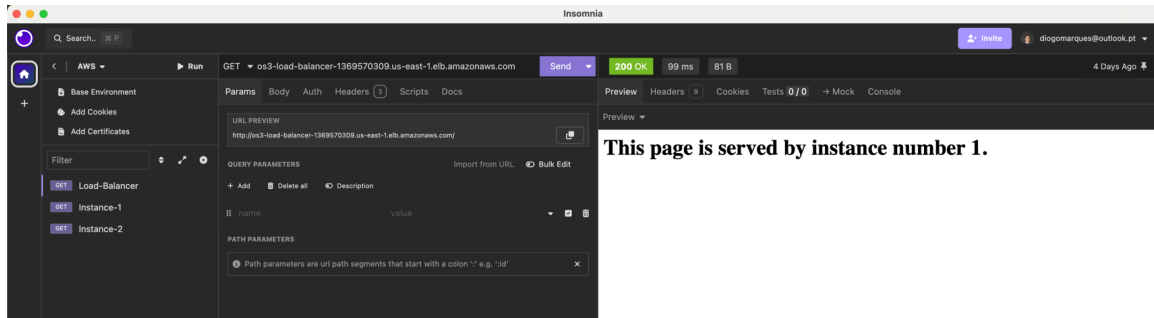
## Add load balancer to network security group



load-balancer-security-group.png

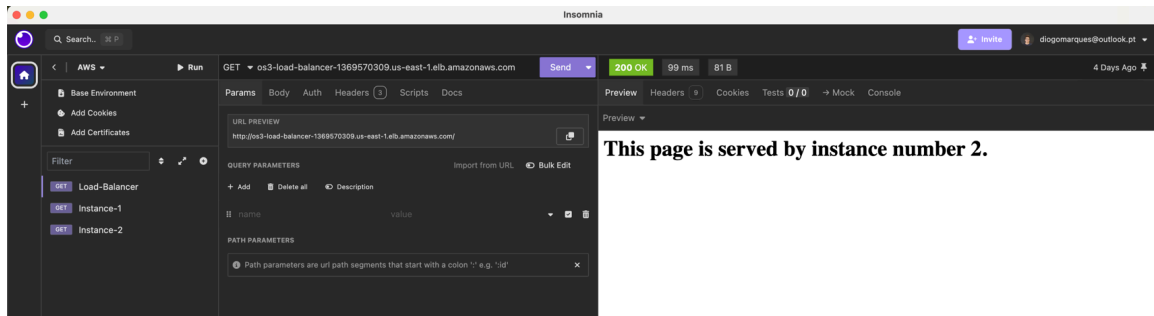
## 8 Test load balancer with Http request to DNS

### Instance 1



This page is served by instance number 1.

### Instance 2



This page is served by instance number 2.

## 9 - 11 50 Request results (measure the time)

Asynchronous Response Time Metrics for 50 Runs:

Average: 0.1016 seconds

Median: 0.0989 seconds

Minimum: 0.0966 seconds

Maximum: 0.2206 seconds

Standard Deviation: 0.0170 seconds

**Metrics:** - Average Response Time: 0.1016 seconds - Median: 0.0989 seconds - Minimum: 0.0966 seconds - Maximum: 0.2206 seconds Standard Deviation: 0.0170 seconds

### Observations:

- The response times are similar to the scenario without the load balancer but show slightly more variability (higher standard deviation).
- Batch 1 has a significantly higher response time (0.2206 seconds), likely due to load balancer initialization or a cold start.
- Subsequent batches stabilize, with times clustering around 0.098-0.100 seconds, similar to the results without the load balancer.

### Comparison

Performance Consistency:

Without the load balancer, the system exhibits slightly more consistent response times (lower standard deviation). With the load balancer, there is marginally higher variability, likely due to the overhead introduced by the load balancer or differences in how requests are distributed to backend servers.



**Initial Latency:**

The first batch with the load balancer exhibits a noticeable delay, suggesting some additional processing or initialization overhead when the load balancer starts handling requests.

**Distribution:**

The similarity in average and median response times suggests that the load balancer effectively distributes requests across backend servers. There is no significant evidence of imbalanced load distribution or server bottlenecks, indicating that the load balancer is performing its function adequately.

**Hypothesis Confirmation**

Based on the data, the load balancer likely employs the Round-Robin technique or a similar balanced distribution strategy. This ensures a near-equal workload across servers, resulting in comparable response times. The marginal increase in variability and overhead could be attributed to: Load balancer processing. Slight differences in backend server performance.

**10 How Does Continuous Requests Influence the Load Balancer?**

Continuous requests to the load balancer can:

Increase the CPU and memory usage on the load balancer due to constant traffic routing. Expose inefficiencies in request distribution if the load balancer is not configured correctly (e.g., Round-Robin, Least Connections, or Weighted Round-Robin). Lead to possible queuing or latency issues, depending on the capacity of the backend servers and the load balancer's processing limits.

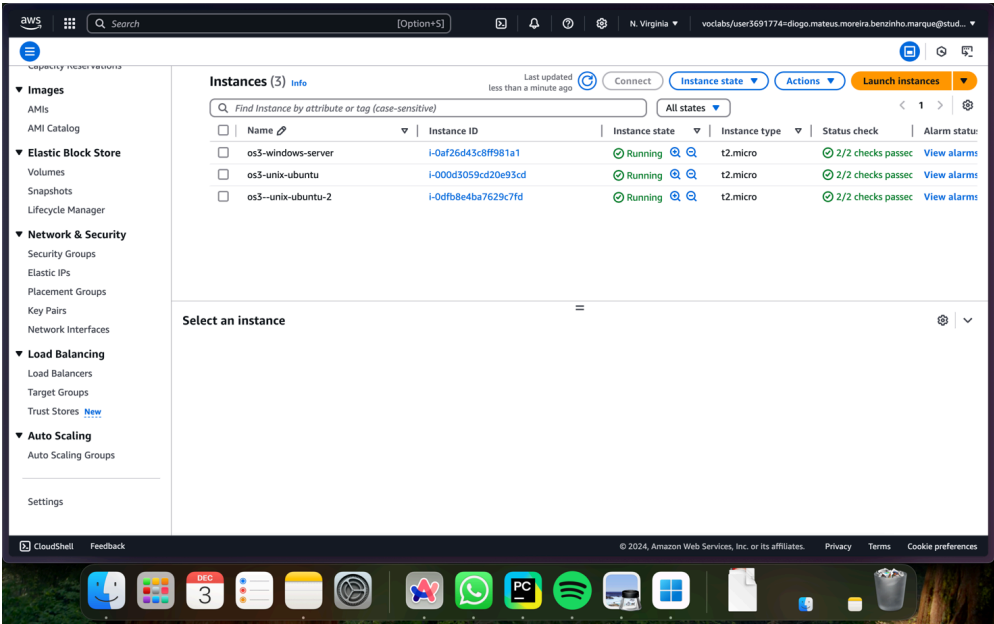
**11 Measure the HTTP response time both on the load balancer as well as on the separate web servers. (9)****12 Conclusion from Results**

From the collected data:

The load balancer and backend servers are handling the load efficiently. The balancing strategy is likely effective.

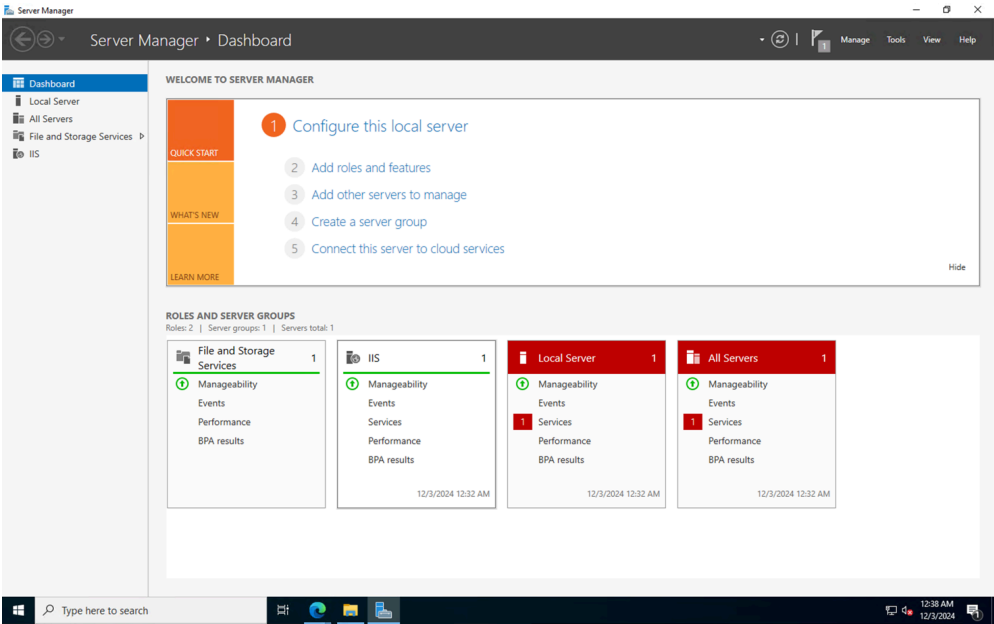
# Windows Server IIS Setup

## Dashboard




dashboard-windows.png

## Setup



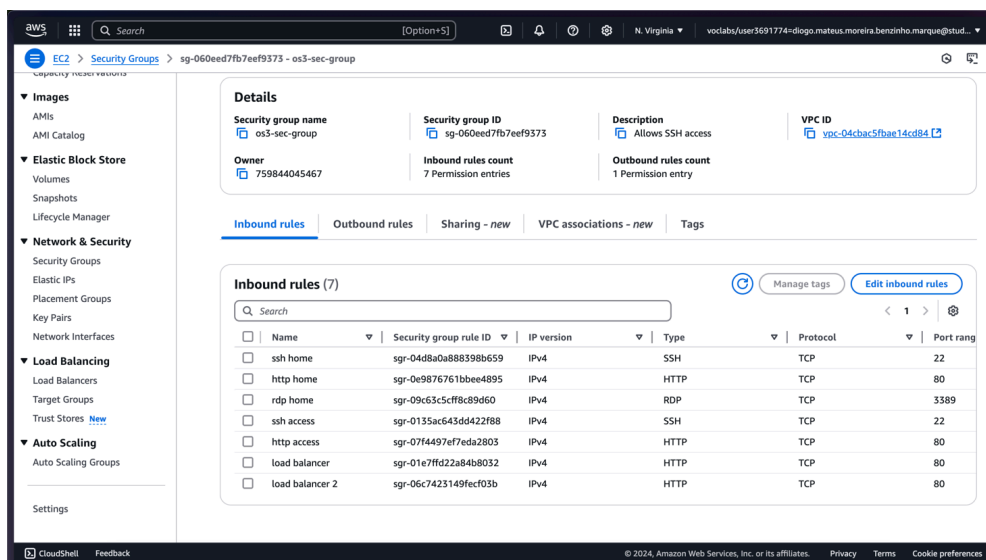
windows-server-iis.png

## HTTP Validation on server



windows-iis.png

## Add to security group to be available on the load balancer



security-group-windows.png

## 13. Advantages/Disadvantages of a Mixed Setup

**Advantages:** - Increased Diversity: Mixing Windows and Ubuntu servers introduces redundancy, reducing downtime risk if one type of server fails.

- **Flexibility:** A mixed environment allows leveraging platform-specific features (e.g., IIS on Windows for .NET applications).
- **Improved Scalability:** Distributing load across different server types can enhance overall system scalability.

**Disadvantages:** - Configuration Complexity: Managing and maintaining a heterogeneous environment can be more challenging.

- **Potential Performance Imbalance:** Differences in server performance may lead to uneven load distribution, requiring weighted load balancing.

- **Software Compatibility:** Additional work may be needed to ensure compatibility across platforms.

## 4 Optional: Security

### 14. Behavioral Change with Three Servers

#### Load balancing including windows server

Asynchronous Response Time Metrics for 50 Runs:

Average: 0.0973 seconds

Median: 0.0953 seconds

Minimum: 0.0935 seconds

Maximum: 0.2010 seconds

Standard Deviation: 0.0149 seconds

#### Response Times with 3 Backend Servers (Ubuntu + IIS)

**Metrics:** - Average Response Time: 0.0973 seconds - Median: 0.0953 seconds - Minimum: 0.0935 seconds - Maximum: 0.2010 seconds Standard Deviation: 0.0149 seconds

**Observations:** - The average response time and median response time are consistent with previous setups involving two Ubuntu servers, indicating that the load balancer effectively distributes requests to the three servers. - The maximum response time (0.2010 seconds) is higher than most other batches, suggesting occasional delays, potentially due to:

- Load balancer overhead. Performance differences between the IIS server and the Ubuntu servers.
- Performance Consistency:
  - The low standard deviation (0.0149 seconds) shows that response times are stable across most batches.
  - The slight variability could stem from differences in how IIS handles requests compared to Ubuntu servers running a similar page.

### Advantages of Adding a Third Server

#### Increased Redundancy:

The addition of an IIS server provides another backend server, reducing the risk of downtime if one server fails. The load balancer now has more options to distribute traffic, potentially improving availability and resilience.

#### Improved Load Distribution:

With three servers, the workload is more evenly distributed, reducing the likelihood of overloading any single server.

### Disadvantages of a Mixed Setup

#### Performance Variability:

If the IIS server has different hardware resources or software configurations, it could handle requests more slowly or differently compared to Ubuntu servers. This variability might increase the response time for some batches, as evidenced by the outlier maximum response time.

**Complexity:**

Managing a mixed environment adds complexity in terms of maintenance, monitoring, and ensuring consistent behavior across platforms. IIS might require different tools or settings compared to Ubuntu, which could increase operational overhead.

**Compatibility Issues:**

While serving static content is straightforward, mixed setups can face compatibility issues with dynamic or application-specific features. Behavioral Changes with Three Servers

**Load Balancer Behavior:**

The load balancer successfully included the IIS server in its rotation, as evidenced by consistent response times across batches. The absence of significant performance degradation suggests that the load balancer's algorithm (likely Round-Robin or Least Connections) distributes traffic effectively.

**Impact of IIS Server:**

The IIS server seems to handle requests comparably to the Ubuntu servers for static content, given the similar average and median response times. Occasional higher response times might indicate slight differences in how requests are processed or resource availability. Conclusion

**Effectiveness of the Load Balancer:**

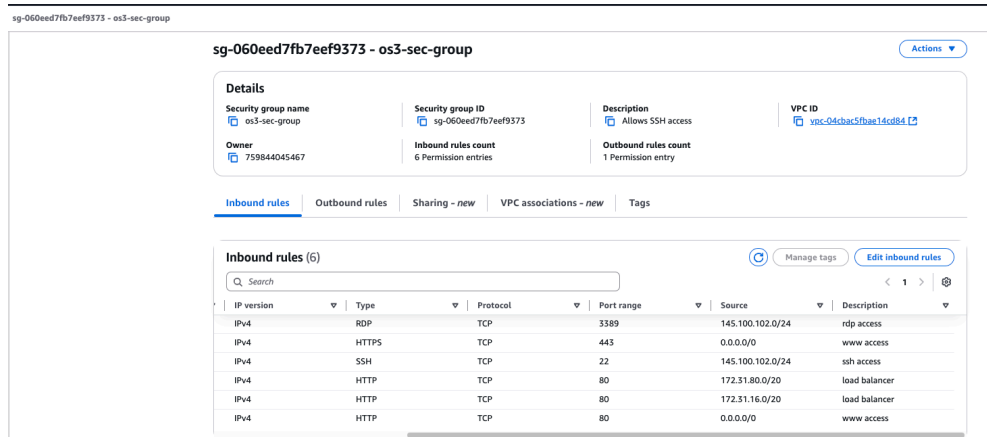
The load balancer continues to distribute requests efficiently, even with three servers. The results suggest that adding the IIS server did not introduce significant issues in handling requests or response time consistency.

**Recommendation for Mixed Environments:**

Mixed setups work well for simple static content, but additional performance monitoring is recommended to identify and address variability between server types. For more complex setups, consider using weighted balancing to account for differences in server performance.

## 15 Modify the Security Group you created in Question 2 for your web servers with reasonable inbound traffic rules

Added HTTP, HTTPS/SSL inbound rules worldwide, RDP/SSH restricted to OS3 network and Load balancer IPs for health checks



security\_group

## 5 Budgeting

16 If you were to build these infrastructure for a small-to-medium sized enterprise, how much should the company budget for 1 year for cloud serv

Based on the AWS cost estimate and infrastructure setup, the company should budget approximately \$419.40 USD for the first year for cloud services.

This budget will cover the following components:

- 2 EC2 Micro Instances running Ubuntu 24.04 for hosting web servers.
- 1 EC2 Micro Instance running Microsoft Windows Server 2022 for additional service hosting.
- 1 Application Load Balancer to distribute incoming traffic across the EC2 instances.

### Breakdown of Estimated Annual Cost:

- **Monthly Cost:** \$34.95 USD.
- **Total Annual Cost:** \$419.40 USD.

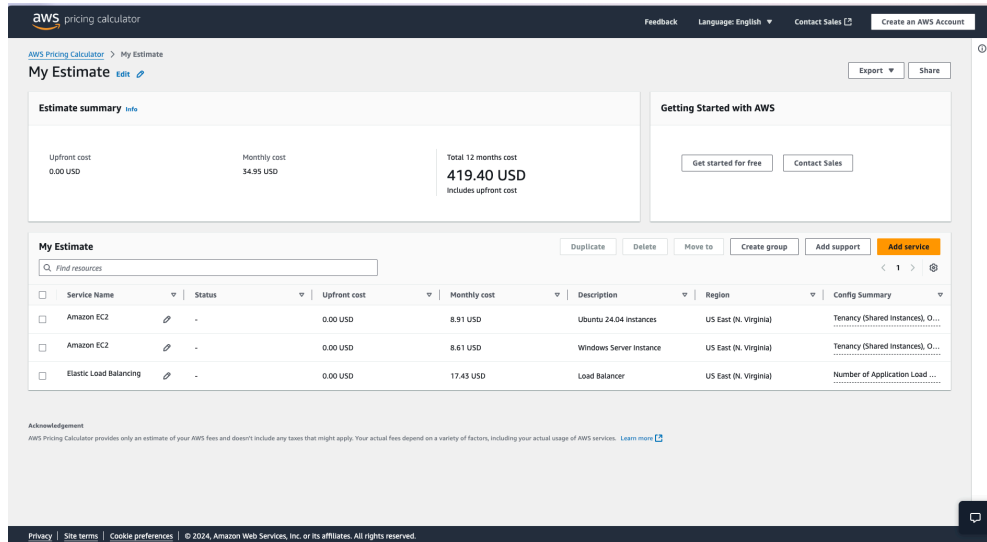
This budget is calculated based on on-demand pricing with no upfront costs and assumes consistent usage throughout the year.

### Considerations:

- **Scalability:** The enterprise may need to adjust the budget if more instances or additional resources (e.g., storage, backup, or data transfer) are needed.
- **Flexibility:** On-demand pricing offers flexibility but is generally higher than reserved instances, so if the company's usage becomes more predictable, they could consider

reserved instances or savings plans in the future to reduce costs. This estimate provides a reasonable starting point for a small-to-medium-sized enterprise running standard web infrastructure on AWS.

**17 AWS Pricing Calculator can generate a persistent link for you to refer to in the future. Provide your estimate's public link for possible discussions with your TAs.**



estimation-costs.png

**Public link to AWS Budget:**

## 6 Virtual Data Center

**18 Do you think this is feasible with the current AWS services?**

Yes, migrating all the infrastructure—storage, networking, servers (web, application, database), and security services—to the Amazon Cloud is not only feasible but also a common approach for many businesses. AWS offers a wide range of services that support these needs, allowing enterprises to move their entire infrastructure to the cloud.

**19 Briefly explain how you would use each AWS service to implement the sysadmin's plan.**

To implement the plan of migrating all infrastructure (storage, network, servers, and security) to AWS, here's a brief explanation of how each service can be used:

### 1. Storage

- **Amazon S3 (Simple Storage Service):**  
Use for storing unstructured data, backups, and static content (e.g., images, logs, documents). S3 provides scalable, durable object storage with low latency.
- **Amazon EBS (Elastic Block Store):**  
Use for persistent block storage attached to EC2 instances. EBS is ideal for databases or file systems where low-latency access is crucial.

- **Amazon EFS (Elastic File System):**  
Use for shared file storage across EC2 instances. Ideal for applications requiring a common file system across multiple servers.

## 2. Networking

- **Amazon VPC (Virtual Private Cloud):**  
Create isolated virtual networks within AWS. Use VPC to segment your infrastructure, set up private subnets, configure route tables, and manage network traffic.
- **AWS Direct Connect:**  
Establish a dedicated network connection from your on-premises data center to AWS. Useful for companies that need high throughput or consistent network performance.
- **Amazon Route 53:**  
Use as a DNS service to manage domain names and route internet traffic to various AWS resources (e.g., EC2 instances, load balancers).

## 3. Servers (Web, Application, Database)

- **Amazon EC2 (Elastic Compute Cloud):**  
Launch and manage scalable compute instances for web, application, and database servers. EC2 provides a wide range of instance types suitable for different workloads, and you can adjust resources as needed.
  - **Web Servers:** Use EC2 instances running Apache or Nginx to serve web pages.
  - **Application Servers:** EC2 can also host application servers that run backend services (e.g., Node.js, Python, or Java).
  - **Database Servers:** EC2 instances can be used to deploy and manage your own database systems (e.g., MySQL, PostgreSQL, etc.), or you can use managed database services (see below).
- **AWS Lambda:**  
For serverless computing. Use Lambda for running small pieces of code (functions) triggered by events, without the need to manage servers.

## 4. Databases

- **Amazon RDS (Relational Database Service):**  
Use RDS for managed relational databases like MySQL, PostgreSQL, Oracle, or SQL Server. RDS simplifies database management by automating tasks like backups, patching, and scaling.
- **Amazon Aurora:**  
A high-performance, fully managed relational database compatible with MySQL and PostgreSQL. Aurora is designed for high availability and read scalability.
- **Amazon DynamoDB:**  
Use for scalable, low-latency NoSQL database needs. Ideal for handling large volumes of unstructured data, and it's fully managed with built-in security and scalability.



## 5. Security

- **AWS Identity and Access Management (IAM):**  
Use IAM to manage users, roles, and permissions across your AWS environment. Define who can access what resources and enforce security policies.
- **Security Groups and Network ACLs:**  
Use these to control inbound and outbound traffic to/from your EC2 instances, ensuring that only authorized users and systems can access your web and application servers.
- **AWS WAF (Web Application Firewall):**  
Protect your applications from common web exploits by using AWS WAF to filter and monitor HTTP traffic.
- **AWS Shield:**  
Use Shield to protect your infrastructure from DDoS attacks.
- **AWS CloudTrail:**  
Track and log API calls and user activity within your AWS environment, which helps with auditing, security monitoring, and compliance.

## Conclusion

Migrating all infrastructure to the Amazon Cloud is definitely feasible, and AWS offers comprehensive services for storage, networking, compute, and security needs. By utilizing the services outlined above, a sysadmin can easily build and manage a complete **Virtual Data Center** in AWS that replaces traditional on-premises infrastructure.

- **Storage:** S3, EBS, EFS
- **Networking:** VPC, Route 53, Direct Connect
- **Servers:** EC2, Lambda
- **Databases:** RDS, Aurora, DynamoDB
- **Security:** IAM, Security Groups, WAF, Shield, CloudTrail

This approach provides flexibility, scalability, and robust security, making it ideal for companies looking to modernize their IT infrastructure and fully embrace cloud computing.